


# COMS W4701: Artificial Intelligence

## Lecture 4: Informed Search



Instructor: Tony Dear

\*Lecture materials derived from UC Berkeley's AI course at [ai.berkeley.edu](https://ai.berkeley.edu)

# Announcements

---

- HW 0 due today!
- HW 1 out now, due in two weeks (Sept 27)
- Weekly review sessions
  - Friday 1-2pm, 214 Pupin
  - Friday 4-5pm, 420 Pupin

# Today

---

- Recap uninformed search
- Heuristics
- Greedy search
- A\* search
  - Optimality
- Admissibility and consistency

# Recap: Search

- Search problem:

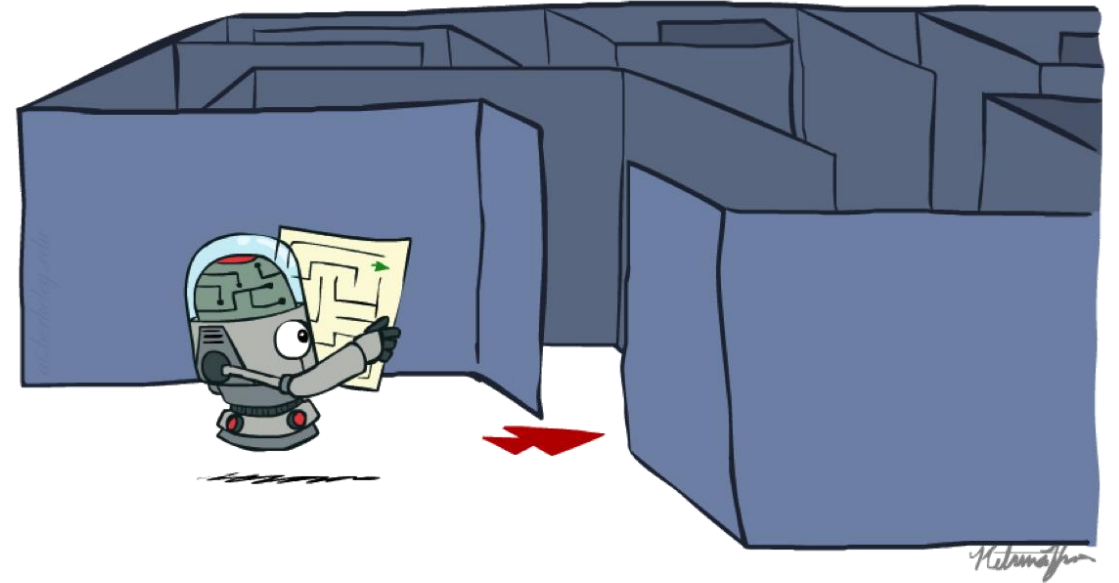
- States (configurations of the world)
- Actions and costs
- Transition model (world dynamics)
- Start state and goal test

- Search tree:

- Nodes: Plans for reaching states
- Plans have costs (sum of action costs)

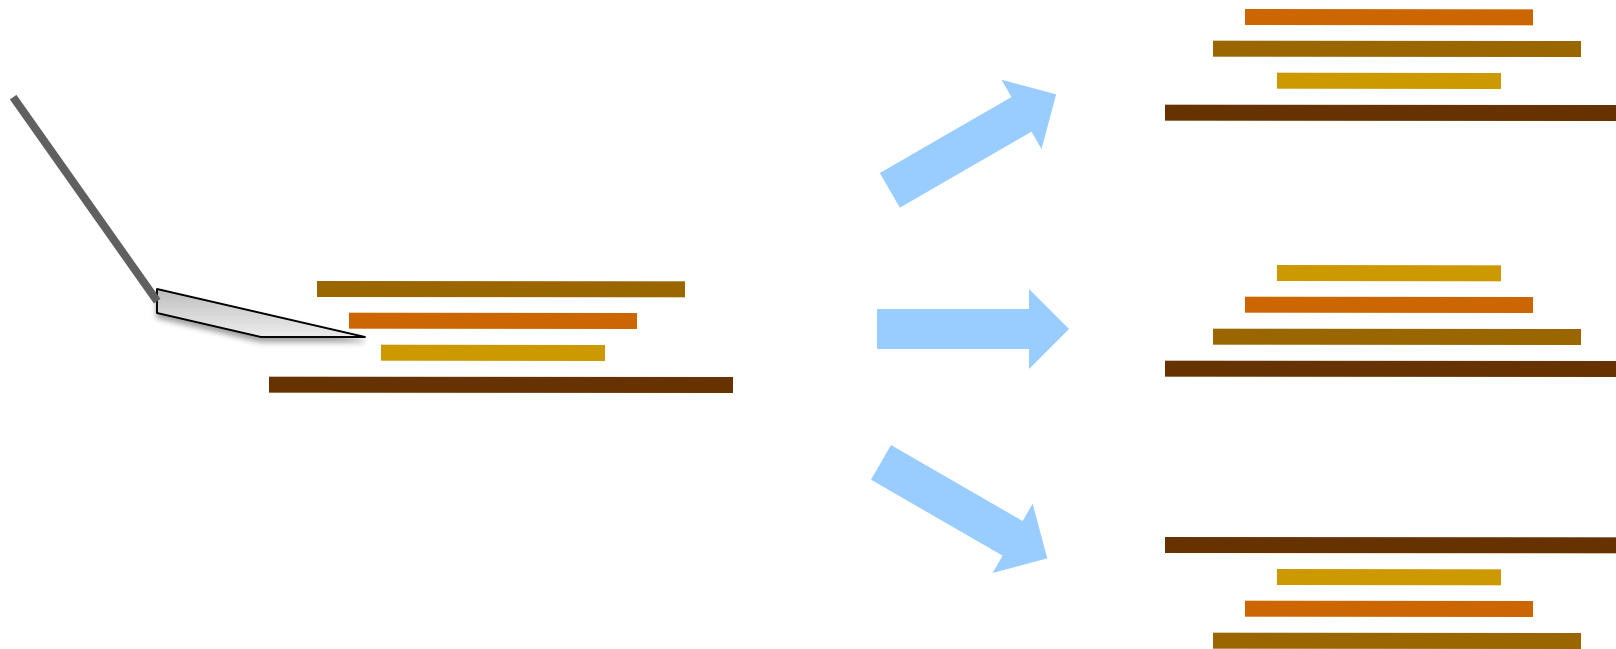
- Search algorithm:

- Systematically builds a search tree
- Chooses an ordering of the fringe (unexplored nodes)
- Optimal: finds least-cost plans



# Example: Pancake Problem


---



Cost: Number of pancakes flipped


# Example: Pancake Problem

State space graph with costs as weights





# General Tree Search

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```



# Redundant States


- Failure to detect repeated states can cause exponentially more work





# Redundant States

- In BFS, for example, we shouldn't bother expanding the circled nodes (why?)




# Graph Search

```
function GRAPH-SEARCH(problem, fringe) return a solution, or failure
  closed ← an empty set
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE[node]) then return node
    if STATE[node] is not in closed then
      add STATE[node] to closed
      for child-node in EXPAND(STATE[node], problem) do
        fringe ← INSERT(child-node, fringe)
      end
  end
```



# Uninformed Search

---





# Uniform Cost Search

- Strategy: expand lowest path cost
- The good: UCS is complete and optimal!
- The bad:
  - Explores options in every “direction”
  - No information about goal location



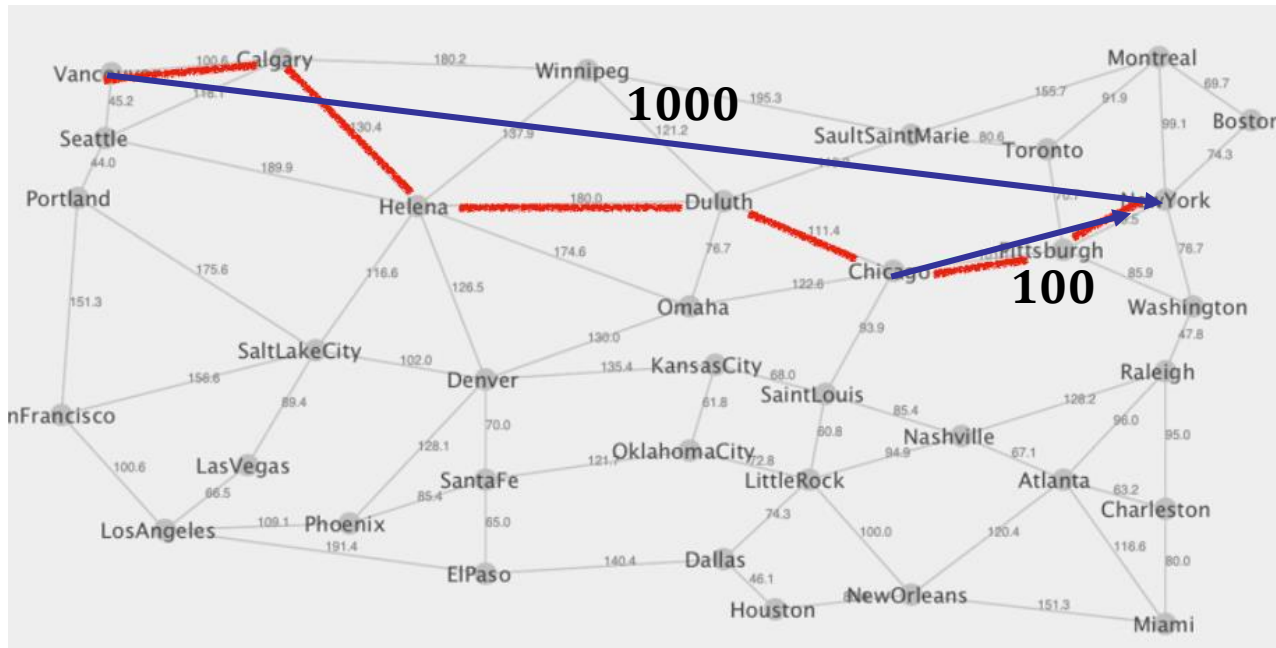


# UCS Examples



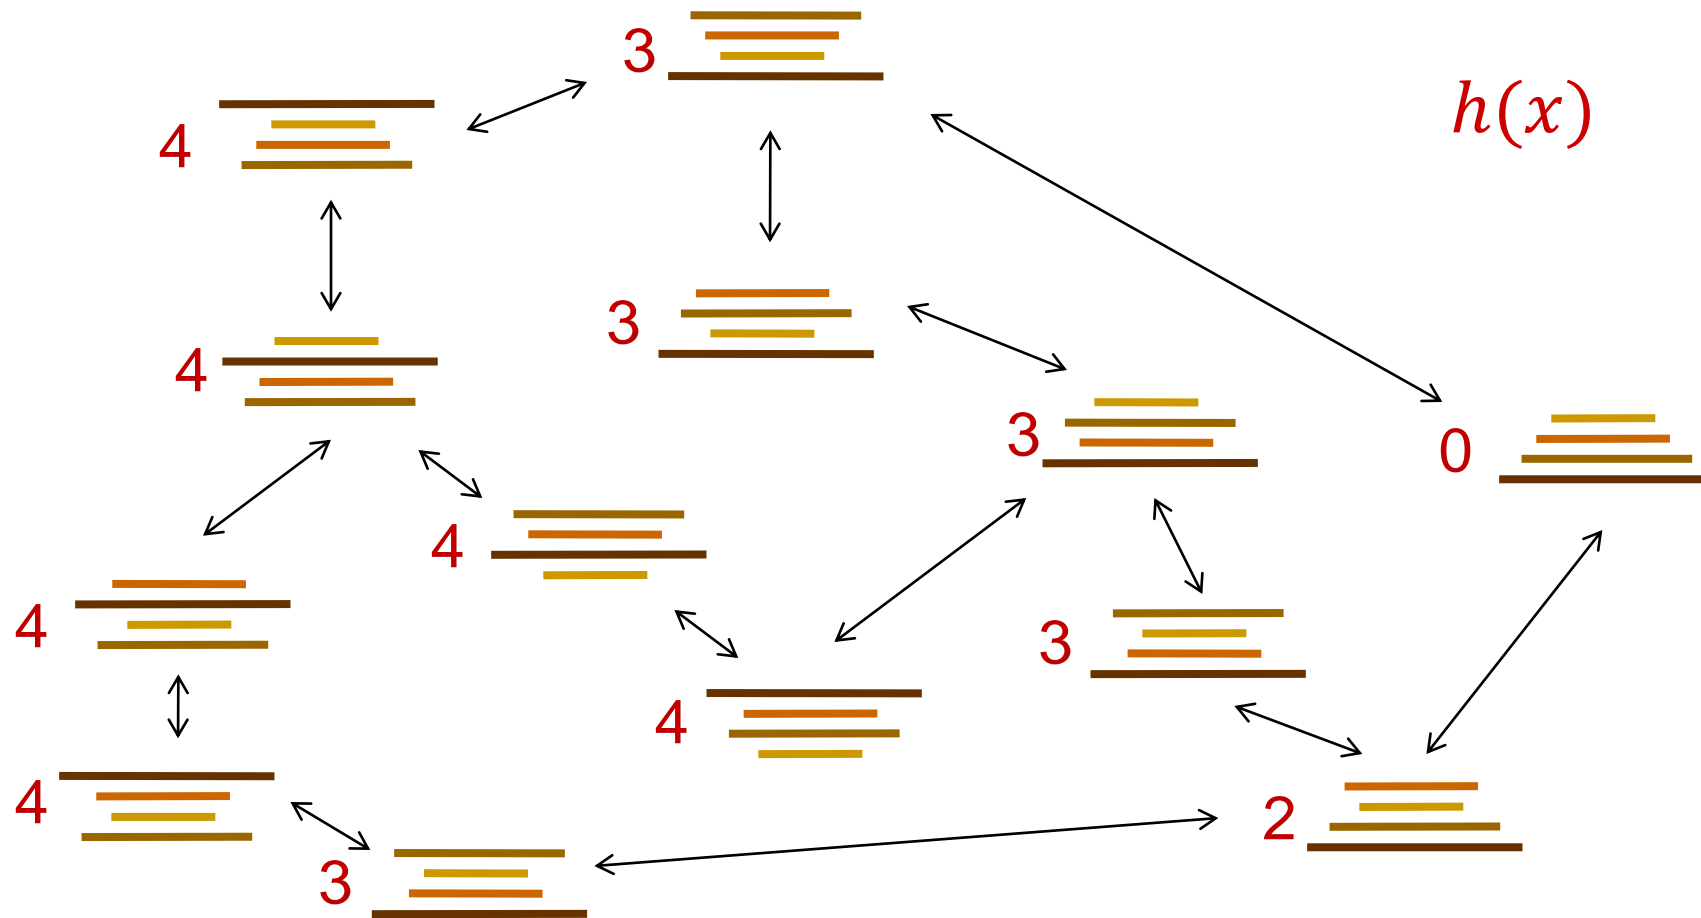
# Search Heuristics

- Heuristic  $h(s)$ 
  - A function that *estimates* how close a state is to a goal
  - Designed for a specific goal
  - Examples: Manhattan distance, Euclidean distance (for pathing)



# Example: Heuristic Function

Heuristic: the number of the largest pancake that is still out of place



# Greedy Search






# Greedy Search

*Strategy: expand node that appears to be closest to the goal*



*Fringe is a priority queue (priority: heuristic)*





- Expand the node that seems closest...
- What can go wrong?

# Greedy Search

- Strategy: expand a node that you think is closest to a goal state
  - Heuristic: estimate of distance to nearest goal for each state
- A common case:
  - Greedy search takes you straight to a goal, regardless of its true cost
- Worst-case: like a badly-guided DFS
  - No guarantee of completeness or optimality!



# Greedy Search Examples




# A\* Search

---



# Combining UCS and Greedy


- **Uniform-cost** orders by path cost, or *backward cost*  $g(n)$
- **Greedy** orders by goal proximity, or *forward cost*  $h(n)$



- **A\* Search** orders by the sum:  $f(n) = g(n) + h(n)$



# When should A\* terminate?

- Should we stop when we enqueue a goal?

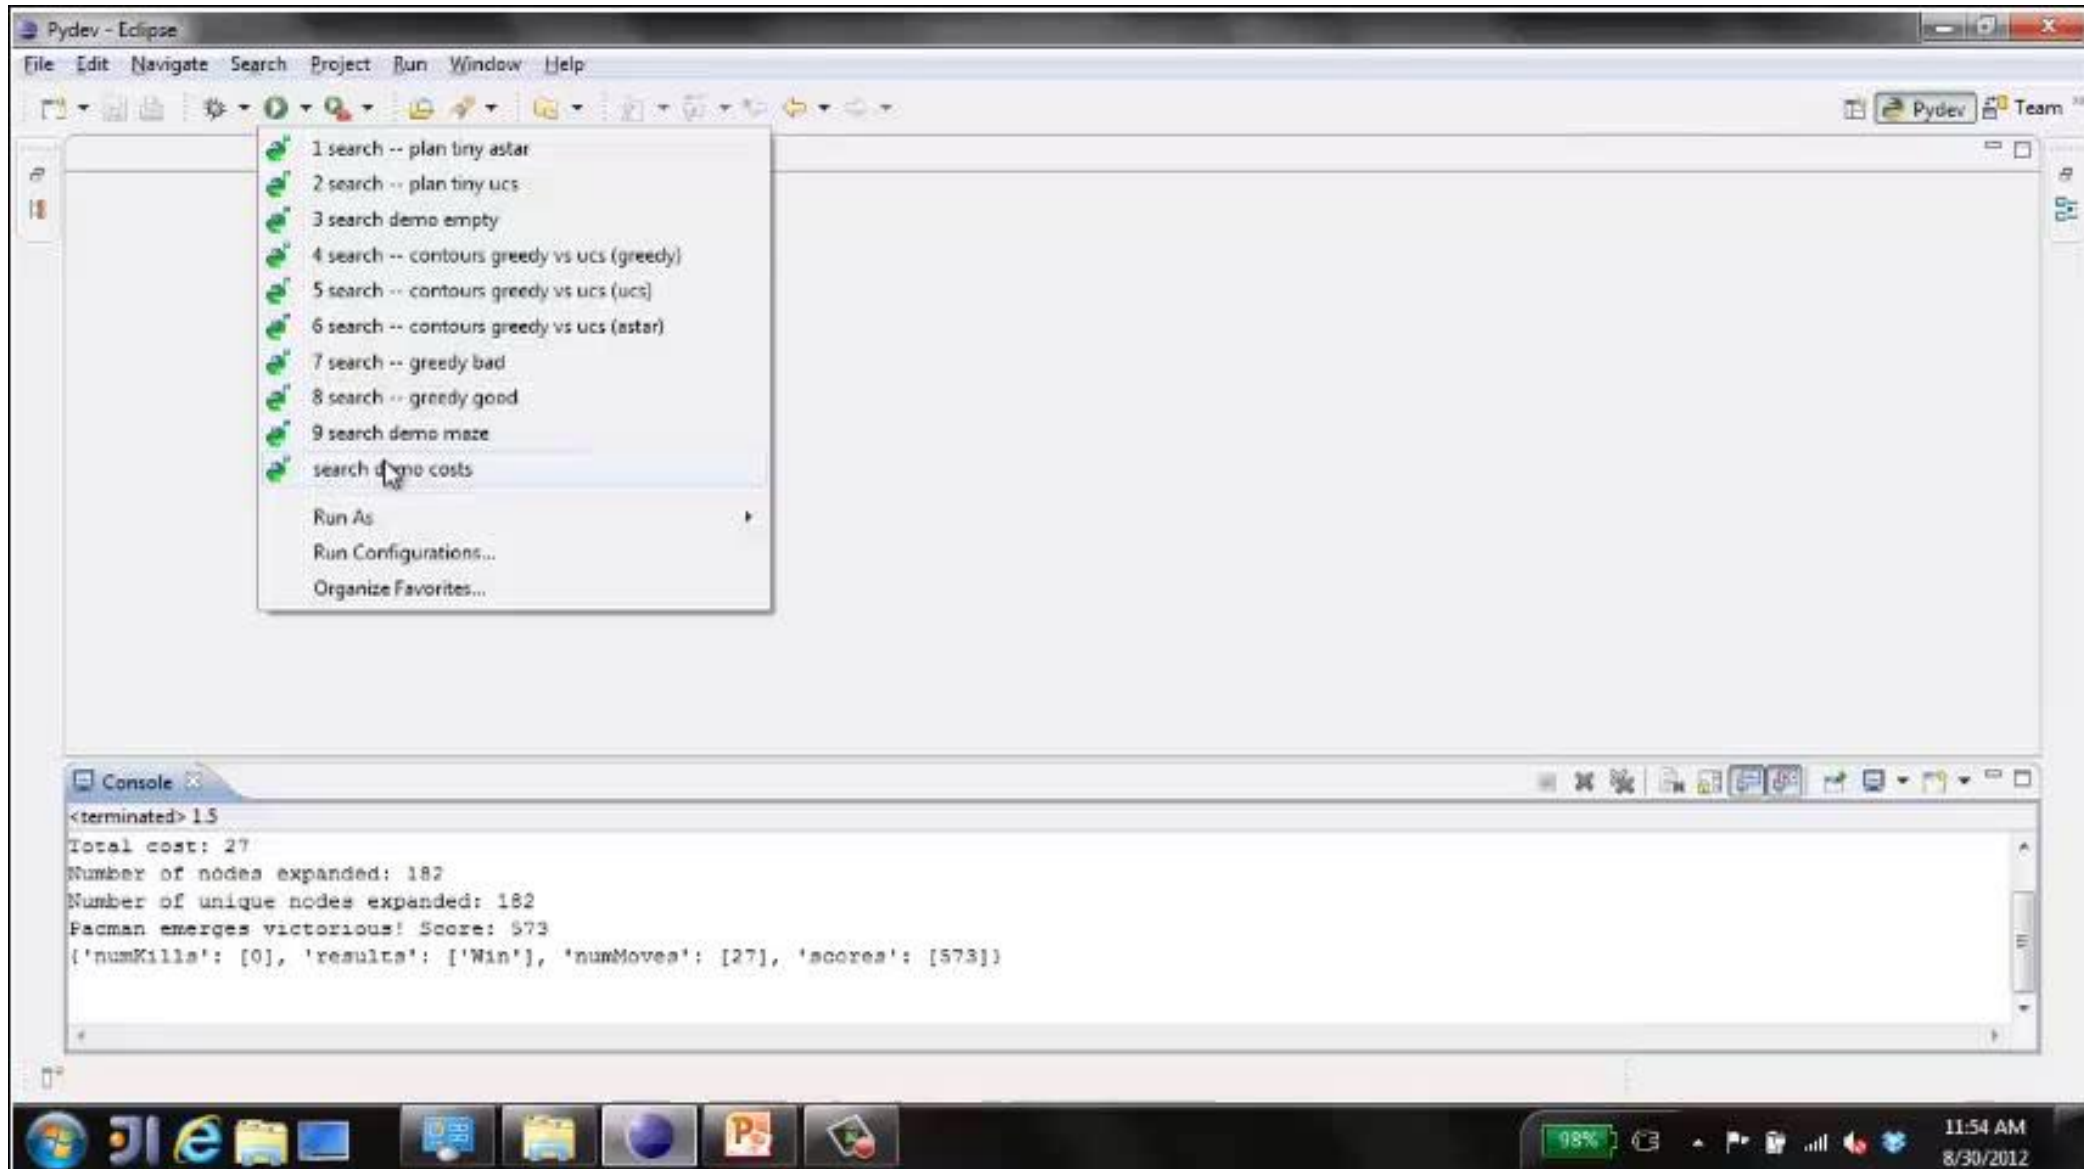


- No: only stop when we dequeue a goal

# A\* Examples



# Guess the Search Algorithm







# A\* Applications

---



- Video games
- Pathing / routing problems
- Resource planning problems
- Robot motion planning
- Language analysis
- Machine translation
- Speech recognition
- ...

# UCS vs A\* Contours


- Uniform-cost expands equally in all “directions”



- A\* expands mainly toward the goal, but does hedge its bets to ensure optimality



# Which solution does A\* return?



S ->

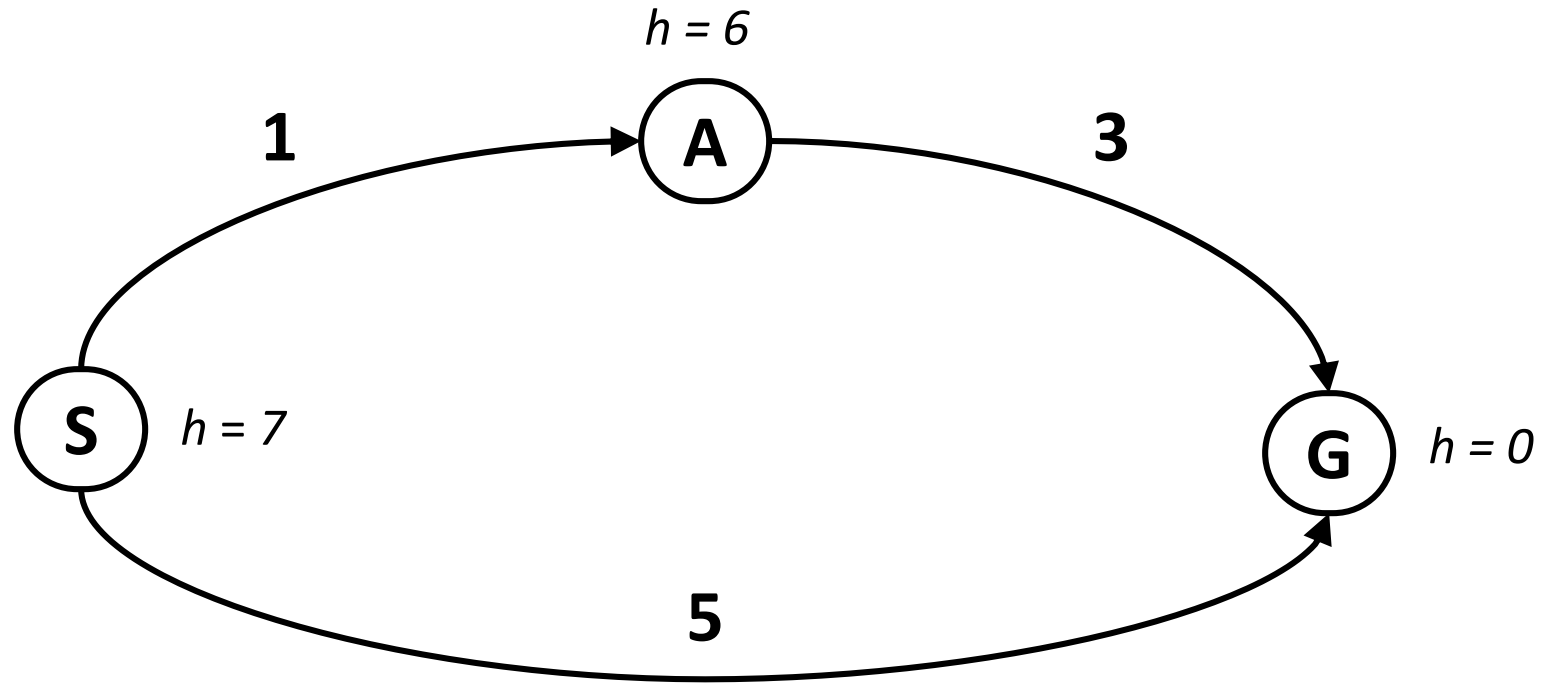
A ->

G

S ->

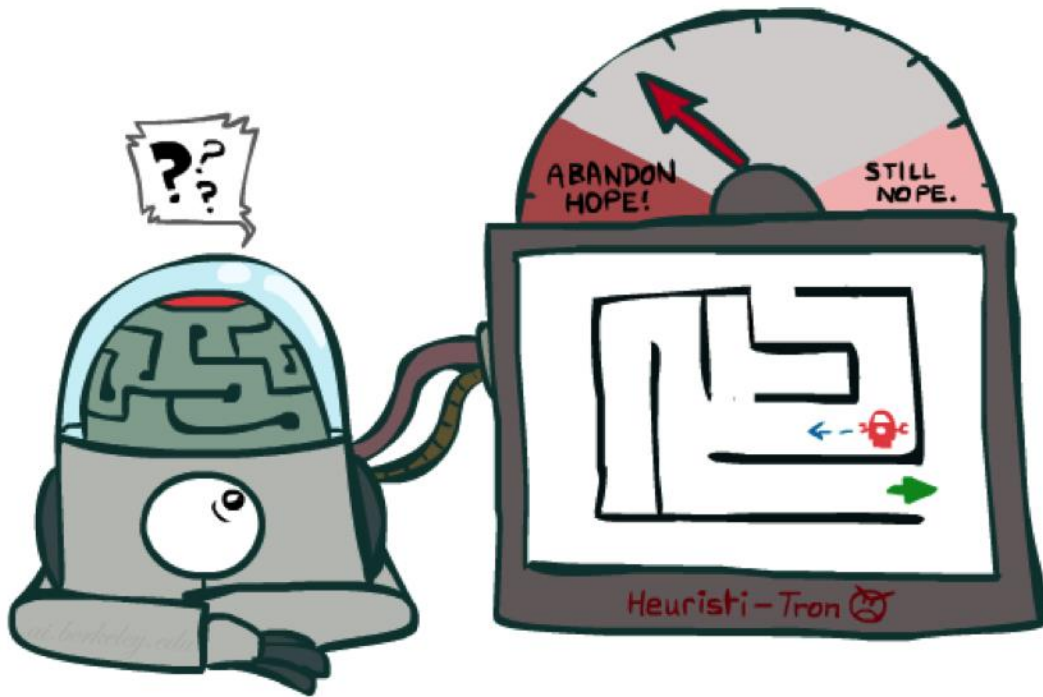
G

# Is A\* Optimal?




- What went wrong?
- Actual bad goal cost < estimated good goal cost
- We need estimates to be less than actual costs!

# Admissible Heuristics



Inadmissible (pessimistic) heuristics break optimality by trapping good plans on the fringe



Admissible (optimistic) heuristics slow down bad plans but never outweigh true costs

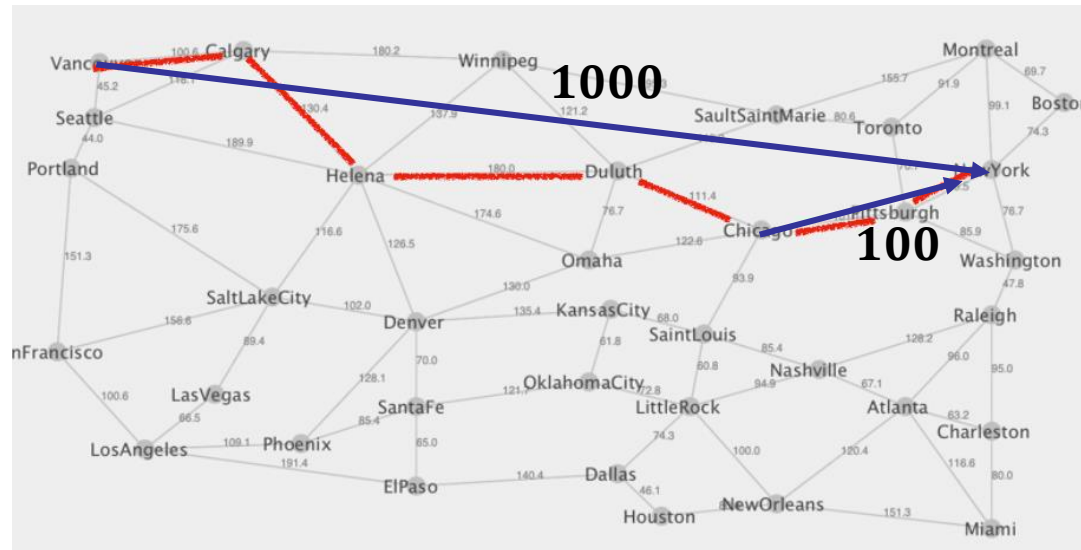
# Admissible Heuristics

- A heuristic  $h$  is *admissible* (optimistic) if:

$$0 \leq h(n) \leq h^*(n)$$

where  $h^*(n)$  is the true cost to a nearest goal

- Examples:




4




# What About the Closed Set?


State space graph



Search tree



# Consistent Heuristics




- Main idea: estimated heuristic costs  $\leq$  actual costs
  - Admissibility: heuristic cost  $\leq$  actual cost to goal
$$h(A) \leq \text{actual cost from A to G}$$
  - Consistency: “heuristic arc” cost  $\leq$  actual cost for each arc
$$h(A) - h(C) \leq \text{cost(A to C)}$$
- What this means:
  - The  $f$  value along a path never decreases
  - True costs  $g$  increase moving toward the goal
  - Estimated costs  $h$  decrease no faster than  $g$  increasing



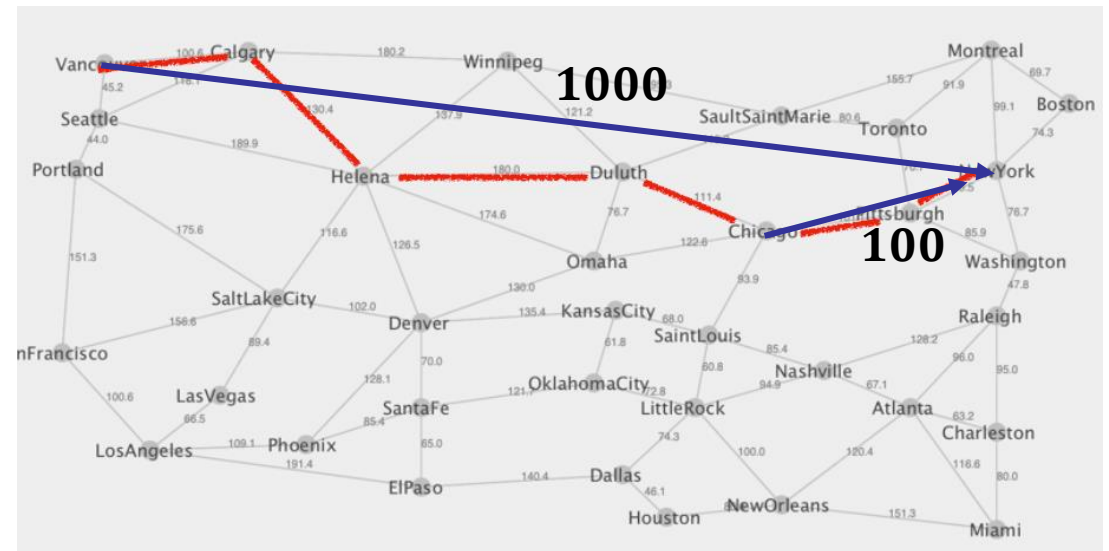
# Optimality of A\* Search

- Sketch: consider what A\* does with a consistent heuristic:
  - Fact 1: A\* expands nodes in increasing total f value (f-contours)
  - Fact 2: Consistency ensures that the first time A\* reaches s, it is along an optimal path
  - Result: A\* graph search is optimal!



# Finding Good Heuristics


- Consistency is a stronger condition—consistency implies admissibility.
- In general, most natural admissible heuristics are also consistent.
- Often, admissible heuristics are solutions to *relaxed problems* with additional actions available.



# Example: 8 Puzzle

7	2	4
5		6
8	3	1

Start State



	1	2
3	4	5
6	7	8

Goal State

# Misplaced Tiles Heuristic

- Heuristic: Number of tiles misplaced
- Why is it admissible?
- This is a *relaxed-problem* heuristic

7	2	4
5		6
8	3	1


Start State

	1	2
3	4	5
6	7	8

Goal State

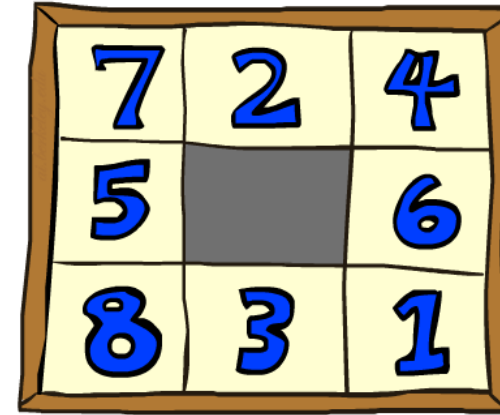
$$h\left(\begin{array}{|c|c|c|}\hline & 1 & 2 \\ \hline 3 & 4 & 5 \\ \hline 6 & 7 & 8 \\ \hline\end{array}\right) = 0$$

$$h\left(\begin{array}{|c|c|c|}\hline 1 & 4 & 2 \\ \hline & 5 & 8 \\ \hline 3 & 6 & 7 \\ \hline\end{array}\right) = 7$$




# Manhattan Distance Heuristic

- What if we had an easier 8-puzzle where any tile could slide any direction at any time, ignoring other tiles?
- Total *Manhattan* distance
- Why is it admissible?
- $h(\text{start}) = 3 + 1 + 2 + \dots = 18$



Start State



Goal State

# Actual Cost?

---

- How about using the *actual cost* as a heuristic?
  - Would it be admissible?
  - Would we save on nodes expanded?
  - What's wrong with it?
- With  $A^*$ : a trade-off between quality of estimate and work per node
  - As heuristics better approximate the true cost, fewer nodes are expanded but more work is needed to compute the heuristic itself

# Summary

---

- Heuristics can guide us toward the goal (à la greedy search)
- $A^*$  uses both backward costs and estimates of forward costs
- $A^*$  is optimal with admissible / consistent heuristics
- Heuristic design often relies on relaxed problems